



**Module Code:** AQ050-3-M-TSF

**Module Name:** Time Series Forecasting

## Individual Assignment

### Japan Military Expenditure Analysis & Forecasting

**Student Name:** Muhammad Yousouf Ali Budullah

**TP Number:** TP086704

**Programme:** MSc Data Science & Business Analytics

**Intake Code:** APDMF2501DSBA(BI)(PR)

**Module Lecturer:** Lee Chee Nian

**Date of Submission:** 19 December 2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Time Series Plot and Component Analysis (a)</b>	<b>2</b>
<b>3</b>	<b>Selected Forecasting Methods (b)</b>	<b>3</b>
<b>4</b>	<b>Train/Test Split and Forecasting (c)</b>	<b>5</b>
<b>5</b>	<b>Forecast Error Evaluation (d)</b>	<b>8</b>
<b>6</b>	<b>Autocorrelation Analysis and Stationarity Testing (e)</b>	<b>9</b>
<b>7</b>	<b>ARIMA/SARIMA Model Development (f)</b>	<b>11</b>
7.1	Proposed Models (i) . . . . .	11
7.2	Model Summary Output (ii) . . . . .	11
7.3	Model Equations (iii) . . . . .	12
7.4	Model Adequacy Checks (iv) . . . . .	12
7.5	Parameter Significance (v) . . . . .	14
7.6	Forecast Errors (vi) . . . . .	15
<b>8</b>	<b>Six-Year Forecast Using Best ARIMA Model (g)</b>	<b>16</b>
<b>A</b>	<b>Appendix: R Code and Additional Outputs</b>	<b>17</b>
	<b>References</b>	<b>32</b>

# **1 Introduction**

For this assignment, Japan was selected as the country for this analysis based on the requirement that the chosen nation must fall within the top 30% of global military expenditure in the most recent year available. Using data from the World Banks Military Expenditure (current US\$) indicator (World Bank, 2025), countries were filtered using R to ensure that they met both criteria of having at least sixty years of available observations and being within the top 30% of expenditure. Japan satisfied these conditions and was therefore chosen for the forecasting study.

The full R code used for the country filtering and data preparation is provided in Appendix A.

## 2 Time Series Plot and Component Analysis (a)

Figure 1 displays the annual military expenditure of Japan from 1960 to 2023. The series demonstrates a strong and persistent upward trend over time, with a sharp increase from the late 1970s through the early 2000s. In the last few decades, the expenditure has fluctuated around a higher value but remains elevated. Although short-term variations are visible, the most dominant feature of the series is its long-term upward trend. Furthermore no repeating seasonal pattern is present, which is to be expected for annual data.




Figure 1: Japan Military Expenditure (1960–2023).

Mann-Kendall test was conducted to confirm the trend observed (see Appendix A for the R code used). The results ( $\tau = 0.855, p < 2.22 \times 10^{-16}$ ) provide strong evidence that the upward trend shown in the plot is statistically significant.

Because the data are recorded annually, seasonality cannot occur (frequency = 1). This was further confirmed using the `isSeasonal()` function (see Appendix A for the R code used), which could not detect seasonality due to the insufficient number of observations per cycle. In addition to the trend, the expenditure displays multi-year rises and falls, indicating the presence of cyclical movements, short-term irregular variations can also be seen.

### 3 Selected Forecasting Methods (b)

Three forecasting methods were selected based on the observed behaviour of the military expenditure series, which shows a clear upward trend and no seasonality. The selected methods are Holts Method, a linear trend regression model, and Simple Exponential Smoothing (SES).

#### Method 1: Holts Method (Double Exponential Smoothing)

Holts Method extends Simple Exponential Smoothing by explicitly modelling both the level and the trend of a time series. The model is defined as follows.

$$\begin{aligned}L_t &= \alpha A_t + (1 - \alpha)(L_{t-1} + T_{t-1}), \\T_t &= \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}, \\\hat{y}_{t+k} &= L_t + kT_t,\end{aligned}$$

where  $A_t$  is the observed value at time  $t$ ,  $L_t$  is the estimated level,  $T_t$  is the estimated trend,  $\alpha$  and  $\beta$  are smoothing parameters ( $0 \leq \alpha \leq 1$ ,  $0 \leq \beta \leq 1$ ), and  $\hat{y}_{t+k}$  is the forecast  $k$  periods ahead.

Holts Method is selected because the military expenditure series is a univariate time series with a clear and persistent upward trend. Hyndman and Athanasopoulos (2021) explain that Holts Method is designed for series with trends and provides suitable forecasts by updating both level and slope over time. In addition, evidence from the M4 Forecasting Competition explains that traditional statistical forecasting methods, particularly those from the exponential smoothing family, consistently perform very well across a wide range of real-world datasets and often outperform more complex approaches (Makridakis et al., 2018). Using Holts Method provides a reliable forecast that avoids the risks of overfitting while accurately capturing the underlying trend of the series.

#### Method 2: Trend Regression Model

Trend regression models capture the long-term relationship between a series and time. A linear trend model is specified as follows.

$$\hat{y}_t = a + bt$$

where  $a$  is the intercept,  $b$  represents the average change in expenditure over time, and  $\hat{y}_t$  is the fitted value at time  $t$ .

A linear trend regression model is included to capture the clear trend observed in the military expenditure series. Hyndman and Athanasopoulos (2021) describe trend regression as a simple yet effective method for modeling the relationship between a forecast variable and a single predictor. Research by Green and Armstrong (2015) further highlights that increasing the complexity of a model does not necessarily improve the accuracy of forecasting and can, in some cases, lead to more errors. Trend regression is used to model the clear and persistent upward trend in expenditure without overreacting to temporary fluctuations.

### **Method 3: Simple Exponential Smoothing (SES)**

Simple Exponential Smoothing produces forecasts by applying exponentially decreasing weights to past observations.

$$F_{t+1} = \alpha A_t + (1 - \alpha)F_t$$

where  $F_{t+1}$  is the forecast for the next period,  $A_t$  is the observed value at time  $t$ , and  $\alpha$  is the smoothing parameter ( $0 \leq \alpha \leq 1$ ).

Forecasting textbooks describe SES as an appropriate method for time series with no clear trend or seasonality (Hyndman & Athanasopoulos, 2021). Although the military expenditure series displays an upward trend, SES is included as a benchmark forecasting method. Petropoulos et al. (2022) highlights the importance of benchmark models in evaluating forecast performance, with SES considered a standard base method in applied forecasting . Including SES provides a benchmark to assess the performance gains of trend-based models like Holts and trend regression.

## 4 Train/Test Split and Forecasting (c)

The full time series data consists of 64 annual observations (1960–2023). A 70/30 split was applied, producing the following.

- Training set: 1960–2004 (45 observations)
- Testing set: 2005–2023 (19 observations)

Forecasts were derived using the methods justified in the previous section, which include Simple Exponential Smoothing (SES), Holts Method, and the linear trend regression model. Each model was fitted using the training data, and forecasts were produced for the testing period. The R code used to perform the train/test split, model estimation, and forecast generation is provided in Appendix A.

The figures below display the observed values, fitted values for the training period, and forecasts for the testing period.

### Observed, Fitted, and Forecast Values




Figure 2: Holts Method: Actual, fitted (training) and forecast (test).

**Linear Trend Model: Actual, Fitted, and Forecast (Japan)**




Figure 3: Linear Trend Model: Actual, fitted (training) and forecast (test).

**SES: Actual, Fitted, and Forecast (Japan)**




Figure 4: Simple Exponential Smoothing (SES): Actual, fitted (training) and forecast (test).

## Estimated Model Parameters

Table 1, Table 2, and Table 3 display the estimated parameters obtained from the fitted forecasting models using the training data (see Appendix A for the R code & output).

Parameter	Estimate
Level smoothing ( $\alpha$ )	0.9900
Initial level ( $l$ )	0.4806

Table 1: Estimated parameters for Simple Exponential Smoothing (SES).

Parameter	Estimate
Level smoothing ( $\alpha$ )	0.9900
Trend smoothing ( $\beta$ )	0.0973
Initial level ( $l$ )	0.4806
Initial trend ( $b$ )	0.0118

Table 2: Estimated parameters for Holts Method.

Coefficient	Estimate	p-value
Intercept	-10.4475	< 0.001
Trend	1.2621	< 0.001

Table 3: Estimated coefficients for the linear trend regression model.

## 5 Forecast Error Evaluation (d)

Mean Absolute Error (MAE) and the Root Mean Squared Error (RMSE) are chosen because they measure forecast errors in the same units as the original series, making the results easy to understand and compare across models. MAE shows the average size of the forecast errors, while RMSE gives more importance to larger errors. Using both measures provides a clear and balanced view of forecast accuracy for this time series.

Forecast accuracy was evaluated using the testing period from 2005 to 2023. All error values are presented in billions of US dollars, consistent with the rescaled time series used during the forecasting stage. The R code used to compute these forecast errors are provided in Appendix A.

Method	RMSE (Billion USD)	MAE (Billion USD)
Simple Exponential Smoothing (SES)	8.15	6.53
Holts Method	6.95	6.05
Linear Trend Regression	12.61	10.66

Table 4: Comparison of forecast errors using RMSE and MAE on the test set (2005–2023).

Based on RMSE and MAE, the Holt Method provides the most accurate forecasts, exhibiting the lowest overall error levels. The linear trend regression model performs the worst, suggesting that a simple deterministic trend is ineffective to capture the behaviour of Japans military expenditure. Simple Exponential Smoothing performs better than the linear model, but falls short of Holts Method, indicating the importance of modeling the underlying upward trend. Consequently, Holts Method is identified as the most appropriate forecasting approach for the Japan Military Expenditure series.

## 6 Autocorrelation Analysis and Stationarity Testing (e)

### Autocorrelation Analysis

Figure 5 displays the ACF and PACF of the original Japan military expenditure series. The ACF exhibits very high autocorrelations at low lags that die down slowly while the PACF shows a single significant spike at lag 1, after which it cuts off and remains close to zero. This pattern is consistent with a non-stationary series.




Figure 5: ACF and PACF of Japan Military Expenditure (Original Series).

### Stationarity Testing

Stationarity was determined using the Augmented Dickey–Fuller (ADF) and Kwiatkowski–Phillips–Schmidt–Shin (KPSS) tests at the 5% significance level. The hypotheses for the tests for ADF and KPSS are defined as follows.

- **ADF test:**

- $H_0$ : The series contains a unit root (non-stationary).
- $H_1$ : The series does not contain a unit root (stationary).

- **KPSS test:**

- $H_0$ : The series is stationary.
- $H_1$ : The series is non-stationary.

## Original Series

Test	Statistic	p-value	Decision (5%)
ADF	-1.5326	0.7643	Unit root not rejected
KPSS	1.5868	0.01	Stationarity rejected

Table 5: Stationarity tests for the original series.

At the 5% significance level, the ADF test does not reject the null hypothesis of a unit root, while the KPSS test rejects the null hypothesis of stationarity. Both tests verify that the series is non-stationary.

## Transformation to Stationarity

To remove non-stationarity, the series was first differenced. The results of the stationarity test for the differenced series are shown in Table 6.

Test	Statistic	p-value	Decision (5%)
ADF	-4.2326	0.01	Unit root rejected
KPSS	0.0984	0.10	Stationarity not rejected

Table 6: Stationarity tests for the differenced series.

For the differenced series, the ADF test rejects the null hypothesis of a unit root, and the KPSS test does not reject the null hypothesis of stationarity. Together, these results confirm that the first-differenced series is stationary.

## 7 ARIMA/SARIMA Model Development (f)

### 7.1 Proposed Models (i)

Figure 6 shows the ACF and PACF of the first-differenced Japan military expenditure series.




Figure 6: ACF and PACF of the differenced series.

Based on this and the results from `auto.arima`, the following models are suggested.

- **ARIMA(1,1,1):** Based on the ACF and PACF of the differenced series, both show a significant spike at lag 1. The ACF appears to cut off after the first lag, while the PACF does not show a clear cut off or dies down pattern at higher lags. This suggests the presence of both autoregressive and moving average components of low order. As a result, an ARIMA(1,1,1) model is proposed.
- **ARIMA(0,1,1) with drift:** In addition, ARIMA(0,1,1) with drift is also suggested as a model, through automatic model suggestion using `auto.arima()`.

### 7.2 Model Summary Output (ii)

Figure 7 presents the summary outputs for the two suggested ARIMA models fitted to Japans military expenditure series: ARIMA(1, 1, 1) and ARIMA(0, 1, 1) with drift.

```
> summary(mod_arima111)
Series: japan_ts
ARIMA(1,1,1)

Coefficients:
          ar1      ma1
        0.2202  0.0911
  s.e.  0.2721  0.2630

sigma^2 = 9.657: log likelihood = -159.86
AIC=325.71  AICc=326.12  BIC=332.14

Training set error measures:
      ME     RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 0.571058 3.033843 2.076269 4.888317 8.601306 0.8906667 -0.02296026

> summary(mod_arima011)
Series: japan_ts
ARIMA(0,1,1) with drift

Coefficients:
          ma1      drift
        0.2490  0.8004
  s.e.  0.1096  0.4714

sigma^2 = 9.326: log likelihood = -158.74
AIC=323.49  AICc=323.89  BIC=329.92

Training set error measures:
      ME     RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 0.002205951 2.981496 2.058092 -11.14334 17.36564 0.8828691 0.02524989
```

ARIMA(1, 1, 1) summary output

ARIMA(0, 1, 1) with drift summary output

Figure 7: Summary outputs for the fitted ARIMA models.

### 7.3 Model Equations (iii)

**ARIMA(1,1,1):**

$$\begin{aligned}(1 - \phi_1 B)(1 - B)y_t &= (1 + \theta_1 B)\varepsilon_t \\ (1 - \phi_1 B)(y_t - y_{t-1}) &= \varepsilon_t + \theta_1 \varepsilon_{t-1} \\ (y_t - y_{t-1}) - \phi_1(y_{t-1} - y_{t-2}) &= \varepsilon_t + \theta_1 \varepsilon_{t-1} \\ y_t &= (1 + \phi_1)y_{t-1} - \phi_1 y_{t-2} + \varepsilon_t + \theta_1 \varepsilon_{t-1}\end{aligned}$$

Using the estimated coefficients  $\hat{\phi}_1 = 0.2202$  and  $\hat{\theta}_1 = 0.0911$ , the fitted model is as follows.

$$y_t = 1.2202 y_{t-1} - 0.2202 y_{t-2} + \varepsilon_t + 0.0911 \varepsilon_{t-1}$$

**ARIMA(0,1,1) with Drift:**

$$\begin{aligned}(1 - B)y_t &= (1 + \theta_1 B)\varepsilon_t + c \\ y_t - y_{t-1} &= \varepsilon_t + \theta_1 \varepsilon_{t-1} + c \\ y_t &= y_{t-1} + c + \varepsilon_t + \theta_1 \varepsilon_{t-1}\end{aligned}$$

Using the estimated parameters  $\hat{\theta}_1 = 0.2490$  and  $\hat{c} = 0.8004$ , the fitted model is as follows.

$$y_t = y_{t-1} + 0.8004 + \varepsilon_t + 0.2490 \varepsilon_{t-1}$$

### 7.4 Model Adequacy Checks (iv)

Model adequacy was evaluated using residual diagnostics and the Ljung–Box test at the 5% significance level. The corresponding code and outputs are shown in Appendix A.

The hypotheses for Ljung–Box is defined as follows.

- $H_0$ : Residuals are not autocorrelated.
- $H_1$ : Residuals are autocorrelated.

## ARIMA(1,1,1)




Figure 8: Residual diagnostics for ARIMA(1,1,1).

Model	Q Statistic	df	p-value	Decision (5%)
ARIMA(1,1,1)	13.168	8	0.1062	Adequate

Table 7: Ljung–Box test results for ARIMA(1,1,1) residuals.

The residuals fluctuate around zero with no clear pattern, and the residual ACF shows no significant autocorrelation. The histogram is approximately normally distributed, indicating near-normal residuals. The Ljung–Box p-value of 0.1062 exceeds 0.05, so the model is adequate.

## ARIMA(0,1,1) with Drift




Figure 9: Residual diagnostics for ARIMA(0,1,1) with drift.

Model	Q Statistic	df	p-value	Decision (5%)
ARIMA(0,1,1) with drift	12.549	9	0.1841	Adequate

Table 8: Ljung–Box test results for ARIMA(0,1,1) with drift residuals.

The residuals show random variation around zero, with no significant autocorrelation in the residual ACF. The histogram is again normally distributed. The Ljung–Box p-value of 0.1841 is greater than 0.05, confirming it is adequate.

Both ARIMA(1,1,1) and ARIMA(0,1,1) with drift pass the residual diagnostics and Ljung–Box test. For both models residuals behave like white noise and satisfy model adequacy conditions.

## 7.5 Parameter Significance (v)

The statistical significance of the ARIMA model parameters was assessed using z-tests at the 5% significance level ( $\alpha = 0.05$ ), with the hypotheses defined as follows.

- $H_0$ : Parameter coefficient = 0.
- $H_1$ : Parameter coefficient  $\neq 0$ .

The corresponding R code and outputs are provided in Appendix A.

### **ARIMA(1,1,1)**

For the ARIMA(1,1,1) model, both the AR(1) coefficient ( $p = 0.4184$ ) and the MA(1) coefficient ( $p = 0.7289$ ) are not statistically significant at the 5% level. Therefore, fails to reject the null hypothesis for either parameter, suggesting that both parameters are not significant.

### **ARIMA(0,1,1) with Drift**

For the ARIMA(0,1,1) with drift model, the MA(1) coefficient is statistically significant at the 5% level ( $p = 0.0231$ ). Therefore the null hypothesis is rejected for this parameter.

At  $\alpha = 0.05$ , none of the parameters in the ARIMA(1,1,1) model are statistically significant, whereas the MA(1) term in the ARIMA(0,1,1) with drift model is significant. This provides stronger evidence for using an ARIMA(0,1,1) model with drift.

## **7.6 Forecast Errors (vi)**

Table 9 presents the forecast error measures for the ARIMA models using the same criteria as in Part (d), namely the RMSE and MAE. All values are expressed in billions of USD.

<b>Model</b>	<b>RMSE</b>	<b>MAE</b>
ARIMA(1,1,1)	3.03	2.08
ARIMA(0,1,1) with Drift	2.98	2.06

Table 9: Forecast error measures (RMSE and MAE) for Holts Method and ARIMA models.

Both ARIMA models achieve lower forecast errors however, the ARIMA(0,1,1) model with drift produces the lowest RMSE and MAE overall. Indicating the strongest forecasting performance for Japans military expenditure. It also suggests that including a moving average component and a drift term improves model performance.

## 8 Six-Year Forecast Using Best ARIMA Model (g)

Figure 10 presents the six-year forecast for Japan's military expenditure using the selected ARIMA(0,1,1) model with drift. Table 10 displays the corresponding point forecasts and 80% and 95% prediction intervals for the forecast horizon. The corresponding code and outputs are shown in Appendix A.




Figure 10: Six-Year Forecast of Japan's Military Expenditure Using the ARIMA(0,1,1) with drift

Year	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2024	51.86	47.95	55.78	45.88	57.85
2025	52.66	46.40	58.93	43.09	62.24
2026	53.46	45.52	61.41	41.32	65.61
2027	54.26	44.94	63.59	39.99	68.53
2028	55.06	44.53	65.60	38.96	71.17
2029	55.87	44.26	67.48	38.11	73.62

Table 10: Six-year ahead forecasts and prediction intervals (Billions USD) generated using ARIMA(0,1,1) with drift.

## A Appendix: R Code and Additional Outputs

The following appendix contains the R code and printed outputs used in the analysis. Only supporting material not shown in the main report is included here.

### Code: Country Filtering and Data Preparation

```
1 library(dplyr)
2
3 # -----
4 # Identify countries that meet requirements (Top 30% expenditure,
5 # with at least 60 years of valid observations)
6 # -----
7
8 final_list <- WB_WDI_MS_MIL_XPND_CD %>%
9   rename(
10     country = REF_AREA_LABEL ,
11     year      = TIME_PERIOD ,
12     mil_exp   = OBS_VALUE
13   ) %>%
14   mutate(year = as.numeric(year)) %>%
15   group_by(country) %>%
16   filter(sum(!is.na(mil_exp)) >= 60) %>%
17   ungroup() %>%
18   filter(year == max(year, na.rm = TRUE)) %>%
19   filter(mil_exp >= quantile(mil_exp, 0.70, na.rm = TRUE)) %>%
20   pull(country)
21
22 final_list
```

### Output: Filtered Countries

```
1 [1] "Japan"
```

*Note.* The filtering procedure returned multiple regions and countries. Only Japan is shown here, as it is the selected country for analysis.

## Code: Extracting and Cleaning Japan's Data

```
1 # Keep relevant variables and rename them
2 df_clean <- WB_WDI_MS_MIL_XPND_CD |>
3   dplyr::select(
4     country = REF_AREA_LABEL ,
5     year    = TIME_PERIOD ,
6     mil_exp = OBS_VALUE
7   ) |>
8   dplyr::mutate(year = as.numeric(year))
9
10 # Extract Japan's data and sort by year
11 japan_military_spending <- df_clean |>
12   dplyr::filter(country == "Japan") |>
13   dplyr::arrange(year)
14
15 str(japan_military_spending)
```

## Output: Structure

```
1 'data.frame': 64 obs. of 3 variables:
2 $ country: chr  "Japan" "Japan" "Japan" "Japan" ...
3 $ year   : num  1960 1961 1962 1963 1964 ...
4 $ mil_exp: num  480555556 492361111 538194444 629861111 718055556
5 ...
```

## Code: Time Series Construction and Trend Testing

```
1 # -----
2 # PART A: Time Series Analysis
3 # -----
4 library(dplyr)
5 library(tseries)
6 library(forecast)
7 library(Kendall)      # Mann Kendall test
8
9 # Convert the dataset into time series (frequency = 1, annual)
10 japan_ts <- ts(japan_military_spending$mil_exp,
11                  start = min(japan_military_spending$year),
12                  frequency = 1)
13
14 # Rescale for plotting
15 japan_ts_bil <- japan_ts / 1e9
16
17 # Time Series Plot
18 plot(japan_ts_bil,
19       main = "Japan Military Expenditure (1960-2023)",
20       xlab = "Year",
21       ylab = "Military Expenditure (Billions USD)",
22       col = "blue",
23       lwd = 2)
24
25 # Mann Kendall Trend Test
26 mk_test <- MannKendall(japan_ts)
27 mk_test
28
29 # Seasonality Test
30 isSeasonal(japan_ts)
```

## Output: Mann-Kendall Test

```
1 tau = 0.855, 2-sided pvalue = < 0.000000000000000222
```

## Output: Seasonality Test

```
1 Error in combined_test(x, freq = freq) :
2   The number of observations per cycle (usually years) is 1 and thus
     too small.
```

## Code: Forecasting Methods and Error Measures

```
1 # -----
2 # PART C & D: Forecasting Methods & Errors
3 # All values expressed in BILLIONS of USD
4 # -----
5
6 library(forecast)
7 library(dplyr)
8
9 # -----
10 # Prepare time series + train/test split
11 # -----
12
13 # Convert to billions for readability
14 japan_ts <- ts(japan_military_spending$mil_exp / 1e9,
15                  start = min(japan_military_spending$year),
16                  frequency = 1)
17
18 n <- length(japan_ts)
19 train_size <- floor(0.7 * n)
20
21 train_end_year <- time(japan_ts)[train_size]
22 test_start_year <- time(japan_ts)[train_size + 1]
23
24 train_japan <- window(japan_ts, end = train_end_year)
25 test_japan <- window(japan_ts, start = test_start_year)
26
27 h <- length(test_japan)
28
29 # -----
30 # Method 1 - Simple Exponential Smoothing (SES)
31 # -----
32
33 ses_model <- ses(train_japan, alpha = 0.99, initial = "simple", h =
34 h)
35
36 plot(ses_model,
37       main = "SES: Actual, Fitted, and Forecast (Japan)",
38       ylab = "Military Expenditure (Billions USD)",
39       xlab = "Year",
40       col = "black",
41       lwd = 2)
42 lines(ses_model$fitted, col = "red", lwd = 2)
```

```

43
44 legend("topleft",
45   legend = c("Actual", "SES Fitted", "SES Forecast"),
46   col = c("black", "red", "#2596be"),
47   lty = 1,
48   lwd = 2)
49
50 # -----
51 # Method 2 - Holt's Method
52 # -----
53
54 holt_model <- holt(train_japan, alpha = 0.99, initial = "simple", h
55   = h)
56
57 plot(holt_model,
58   main = "Holt: Actual, Fitted, and Forecast (Japan)",
59   ylab = "Military Expenditure (Billions USD)",
60   xlab = "Year",
61   col = "black",
62   lwd = 2)
63
64 lines(holt_model$fitted, col = "red", lwd = 2)
65
66 legend("topleft",
67   legend = c("Actual", "Holt Fitted", "Holt Forecast"),
68   col = c("black", "red", "#2596be"),
69   lty = 1,
70   lwd = 2)
71
72 # -----
73 # Method 3 - Linear Trend Regression
74 # -----
75
76 linear_model <- tslm(train_japan ~ trend)
77
78 test_trend <- data.frame(
79   trend = seq(length(train_japan) + 1,
80             length(train_japan) + h)
81 )
82
83 linear_forecast <- forecast(linear_model,
84   newdata = test_trend,
85   h = h)
86
87 plot(linear_forecast,

```

```

87 main = "Linear Trend Model: Actual, Fitted, and Forecast (Japan
88           )",
89 ylab = "Military Expenditure (Billions USD)",
90 xlab = "Year")
91
91 lines(fitted(linear_model), col = "red", lwd = 2)
92
93 legend("topleft",
94         legend = c("Actual", "Linear Fitted", "Linear Forecast"),
95         col = c("black", "red", "#2596be"),
96         lty = 1,
97         lwd = 2,
98         cex = 0.8)
99
100 # -----
101 # 5. Display Parameters
102 # -----
103 ses_model$model$par
104 holt_model$model$par
105 summary(linear_model)
106
107 # -----
108 # Forecast Errors (Billions USD)
109 # -----
110
111 ses_acc <- accuracy(ses_model, test_japan)
112 holt_acc <- accuracy(holt_model, test_japan)
113 linear_acc <- accuracy(linear_forecast$mean, test_japan)
114
115 ses_acc
116 holt_acc
117 linear_acc

```

## Output: Parameters

```

1 > ses_model$model$par
2   alpha      1
3 0.9900000 0.4805556
4 > holt_model$model$par
5   alpha      beta      1          b
6 0.99000000 0.09732261 0.48055556 0.01180556
7 > summary(linear_model)
8
9 Call:

```

```

10 | tslm(formula = train_japan ~ trend)
11 |
12 Residuals:
13   Min     1Q Median     3Q    Max
14 -9.2789 -4.6099 -0.8813  3.8252 14.9742
15
16 Coefficients:
17             Estimate Std. Error t value      Pr(>|t|)
18 (Intercept) -10.44747   1.84600  -5.66 0.00000123 ***
19 trend        1.26208   0.07145  17.66 < 0.0000000000000002 ***
20 ---
21 Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 .
22          0.1      1
23
24 Residual standard error: 6.018 on 42 degrees of freedom
25 Multiple R-squared:  0.8814, Adjusted R-squared:  0.8785
F-statistic: 312 on 1 and 42 DF, p-value: < 0.0000000000000002

```

## Output: Forecast Error

```

1 > ses_acc
2               ME      RMSE      MAE      MPE      MAPE
3                   MASE      ACF1 Theil's U
4 Training set 0.9635961 2.662398 1.836646  9.237544 11.58833
5           0.9823245 0.3164454      NA
6 Test set     6.2077674 8.149869 6.525333 11.776811 12.55164
7           3.4900536 0.6903032 1.903656
8 > holt_acc
9               ME      RMSE      MAE      MPE      MAPE
10                  MASE      ACF1 Theil's U
11 Training set 0.1962618 2.529633 1.585638  4.906766 9.009683
12           0.8480735 0.2746422      NA
13 Test set     -2.6585572 6.947505 6.050857 -6.371179 12.330608
14           3.2362816 0.7560415 1.733011
15 > linear_acc
16               ME      RMSE      MAE      MPE      MAPE      ACF1
17                   Theil's U
18 Test set -9.673382 12.60658 10.66208 -20.8112 22.45925 0.7951786
19           3.26078

```

## Code: Autocorrelation and Stationarity Analysis

```
1 # -----
2 # PART E: Autocorrelation & Stationarity Analysis
3 # -----
4
5 library(forecast)
6 library(tseries)
7
8 # -----
9 # Time series object (anual data)
10 # -----
11
12 japan_ts <- ts(japan_military_spending$mil_exp,
13                  start = min(japan_military_spending$year),
14                  frequency = 1)
15
16 # -----
17 # Autocorrelation analysis (original series)
18 # -----
19
20 acf(japan_ts,
21      main = "ACF of Japan Military Expenditure")
22
23 pacf(japan_ts,
24       main = "PACF of Japan Military Expenditure")
25
26 # -----
27 # Stationarity tests (original series)
28 # -----
29
30 # Augmented Dickey Fuller Test
31 adf_test <- adf.test(japan_ts)
32 adf_test
33
34 # KPSS Test
35 kpss_test <- kpss.test(japan_ts)
36 kpss_test
37
38 # Suggested differencing order
39 ndiffs(japan_ts)
40
41 # -----
42 # First differencing
43 # -----
```

```

44
45 japan_diff <- diff(japan_ts)
46 #Convert to billions
47 plot(japan_diff/ 1e9,
48      main = "Differenced Japan Military Expenditure",
49      ylab = "Differenced Values (Billions USD)",
50      xlab = "Year",
51      col = "blue")
52
53 # -----
54 # Stationarity tests (on differenced series)
55 # -----
56
57 adf_diff <- adf.test(japan_diff)
58 adf_diff
59
60 kpss_diff <- kpss.test(japan_diff)
61 kpss_diff
62
63 # -----
64 # Autocorrelation analysis (on differenced series)
65 # -----
66
67 acf(japan_diff,
68      main = "ACF of Differenced Series")
69
70 pacf(japan_diff,
71       main = "PACF of Differenced Series")
72

```

## Output: Stationarity Tests

```

1 > adf_test
2
3 Augmented Dickey-Fuller Test
4
5 data: japan_ts
6 Dickey-Fuller = -1.5326, Lag order = 3, p-value = 0.7643
7 alternative hypothesis: stationary
8
9 > kpss_test
10
11 KPSS Test for Level Stationarity
12
13 data: japan_ts

```

```
14 KPSS Level = 1.5868, Truncation lag parameter = 3, p-value = 0.01
15
16 > ndiffs(japan_ts)
17 [1] 1
18
19 > adf_diff
20
21 Augmented Dickey-Fuller Test
22
23 data: japan_diff
24 Dickey-Fuller = -4.2326, Lag order = 3, p-value = 0.01
25 alternative hypothesis: stationary
26
27 > kpss_diff
28
29 KPSS Test for Level Stationarity
30
31 data: japan_diff
32 KPSS Level = 0.098352, Truncation lag parameter = 3, p-value = 0.1
```

## Code: ARIMA Modelling

```
1 # -----
2 # PART F: ARIMA Modeling
3 # -----
4 library(forecast)
5 library(tseries)
6 library(lmtest)
7
8 # -----
9 # Define time series (scaled to billions USD)
10 # -----
11
12 japan_ts <- ts(japan_military_spending$mil_exp/1e9,
13                   start = min(japan_military_spending$year),
14                   frequency = 1)
15
16 # -----
17 # (f)(i) Select ARIMA models
18 # -----
19 auto.arima(japan_ts, trace = TRUE)
20
21 # -----
22 # (f)(ii) Fit ARIMA models
23 # -----
24 # Values in billions as japan_ts was scaled
25
26 # ARIMA(1,1,1)
27 mod_arima111 <- Arima(japan_ts, order = c(1,1,1))
28
29 # ARIMA(0,1,1) with drift
30 mod_arima011 <- Arima(japan_ts, order = c(0,1,1), include.drift =
31   TRUE)
32
33 summary(mod_arima111)
34 summary(mod_arima011)
35
36 # -----
37 # (f)(iv) Model adequacy checks
38 # -----
39 checkresiduals(mod_arima111)
40 checkresiduals(mod_arima011)
41
42 # -----
43 # (f)(v) Coefficient significance tests
```

```

43 # -----
44 coefstest(mod_arima111)
45 coefstest(mod_arima011)

46 #
47 # -----
48 # (f)(vi) Forecast error measures
49 #
50 accuracy(mod_arima111)
51 accuracy(mod_arima011)

```

## Output: Auto Arima

```

1 > auto.arima(japan_ts, trace = TRUE)
2
3 ARIMA(2,1,2) with drift      : Inf
4 ARIMA(0,1,0) with drift      : 326.0715
5 ARIMA(1,1,0) with drift      : 323.9396
6 ARIMA(0,1,1) with drift      : 323.8947
7 ARIMA(0,1,0)                  : 327.8553
8 ARIMA(1,1,1) with drift      : 325.9804
9 ARIMA(0,1,2) with drift      : 325.1067
10 ARIMA(1,1,2) with drift      : 325.7833
11 ARIMA(0,1,1)                  : 324.4632
12
13 Best model: ARIMA(0,1,1) with drift
14
15 Series: japan_ts
16 ARIMA(0,1,1) with drift
17
18 Coefficients:
19       ma1     drift
20       0.2490  0.8004
21 s.e.  0.1096  0.4714
22
23 sigma^2 = 9.326: log likelihood = -158.74
24 AIC=323.49    AICc=323.89    BIC=329.92

```

## Output: Model Summaries

```

1 > summary(mod_arima111)
2 Series: japan_ts
3 ARIMA(1,1,1)
4
5 Coefficients:

```

```

6      ar1      ma1
7      0.2202  0.0911
8 s.e.  0.2721  0.2630
9
10 sigma^2 = 9.657: log likelihood = -159.86
11 AIC=325.71    AICc=326.12    BIC=332.14
12
13 Training set error measures:
14             ME      RMSE      MAE      MPE      MAPE      MASE
15             ACF1
16 Training set 0.571058 3.033843 2.076269 4.888317 8.601306 0.8906667
17 -0.02296026
18 > summary(mod_arima011)
19 Series: japan_ts
20 ARIMA(0,1,1) with drift
21
22 Coefficients:
23             ma1      drift
24             0.2490  0.8004
25 s.e.  0.1096  0.4714
26
27 sigma^2 = 9.326: log likelihood = -158.74
28 AIC=323.49    AICc=323.89    BIC=329.92
29
30 Training set error measures:
31             ME      RMSE      MAE      MPE      MAPE
32             MASE      ACF1
33 Training set 0.002205951 2.981496 2.058092 -11.14334 17.36564
34 0.8828691 0.02524989

```

## Output: Residual Diagnostics (Ljung–Box Test)

```

1 > checkresiduals(mod_arima111)
2
3 Ljung-Box test
4
5 data: Residuals from ARIMA(1,1,1)
6 Q* = 13.168, df = 8, p-value = 0.1062
7
8 Model df: 2. Total lags used: 10
9
10 > checkresiduals(mod_arima011)
11
12 Ljung-Box test

```

```

13
14 data: Residuals from ARIMA(0,1,1) with drift
15 Q* = 12.549, df = 9, p-value = 0.1841
16
17 Model df: 1. Total lags used: 10

```

## Output: Coefficient Significance Tests

```

1 > coeftest(mod_arima111)
2
3 z test of coefficients:
4
5 Estimate Std. Error z value Pr(>|z|)
6 ar1 0.220158 0.272093 0.8091 0.4184
7 ma1 0.091142 0.262952 0.3466 0.7289
8
9 > coeftest(mod_arima011)
10
11 z test of coefficients:
12
13 Estimate Std. Error z value Pr(>|z|)
14 ma1 0.24895 0.10958 2.2720 0.02309 *
15 drift 0.80039 0.47142 1.6978 0.08954 .
16 ---
17 Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 .
0.1 1

```

## Output: Forecast Accuracy (Training Set)

```

1 > accuracy(mod_arima111)
2
3 ME RMSE MAE MPE MAPE MASE
4 ACF1
5 Training set 0.571058 3.033843 2.076269 4.888317 8.601306 0.8906667
6 -0.02296026
7 > accuracy(mod_arima011)
8
9 ME RMSE MAE MPE MAPE
10 MASE ACF1
11 Training set 0.002205951 2.981496 2.058092 -11.14334 17.36564
12 0.8828691 0.02524989

```

## Six-Year Forecast Using ARIMA(0,1,1) with Drift

```
1 # -----
2 # PART G: Forecast 6 years
3 # -----
4
5 # Forecast the next 6 years
6 forecast_6yr <- forecast(mod_arima011, h = 6)
7
8 # Print forecast values
9 forecast_6yr
10
11 # -----
12 # Plot forecast and fitted values
13 # -----
14
15 plot(forecast_6yr,
16       main = "6-Year Forecast of Japan Military Expenditure\nARIMA
17             (0,1,1) with Drift",
18       ylab = "Military Expenditure (Billions USD)",
19       xlab = "Year",
20       col = "black",
21       lwd = 2)
22
23 # Add fitted values
24 lines(fitted(mod_arima011), col = "red", lwd = 2)
25
26 # Add legend
27 legend("topleft",
28        legend = c("Actual", "ARIMA(0,1,1) Fitted", "ARIMA(0,1,1)
29                      Forecast"),
30        col = c("black", "red", "#2596be"),
31        lty = 1,
32        lwd = 2)
```

## Output: Forecast of next 6 years

```
1 > forecast_6yr
2   Point Forecast    Lo 80     Hi 80    Lo 95     Hi 95
3 2024      51.86341 47.94964 55.77718 45.87781 57.84900
4 2025      52.66379 46.40191 58.92567 43.08707 62.24052
5 2026      53.46418 45.52033 61.40803 41.31511 65.61325
6 2027      54.26456 44.93727 63.59185 39.99970 68.52942
7 2028      55.06495 44.53442 65.59547 38.95990 71.16999
8 2029      55.86533 44.25562 67.47505 38.10980 73.62086
```

## References

- Green, K. C., & Armstrong, J. S. (2015). Simple versus complex forecasting: The evidence. *Journal of Business Research*, 68(8), 1678–1685. <https://doi.org/10.1016/j.jbusres.2015.03.026>
- Hyndman, R. J., & Athanasopoulos, G. (2021). *Forecasting: Principles and practice* (3rd ed.) [Accessed 2025]. OTexts. <https://otexts.com/fpp3/>
- Makridakis, S., Spiliotis, E., & Assimakopoulos, V. (2018). The m4 competition: Results, findings, conclusion and way forward. *International Journal of Forecasting*, 34(4), 802–808. <https://doi.org/10.1016/j.ijforecast.2018.06.001>
- Petropoulos, F., Makridakis, S., Assimakopoulos, V., & Nikolopoulos, K. (2022). Forecasting: Theory and practice. *International Journal of Forecasting*, 38(3), 705–871. <https://doi.org/10.1016/j.ijforecast.2021.11.001>
- World Bank. (2025). Military expenditure (current us\$) [Retrieved from [https://data360.worldbank.org/indicator/WB\\_WDI\\_MS\\_MIL\\_XPND\\_CD](https://data360.worldbank.org/indicator/WB_WDI_MS_MIL_XPND_CD)].